

Transactions, Views, Indexes

CONTROLLING CONCURRENT BEHAVIOR

VIRTUAL AND MATERIALIZED VIEWS

SPEEDING ACCESSES TO DATA

Why Transactions?

Database systems are normally being accessed by many users or processes at the same time.

- Both queries and modifications.

The same as operating systems, which *support* interaction of processes, a DMBS needs to keep processes from troublesome interactions.

Example: Bad Interaction

You and your partner each take \$100 from different ATM's at about the same time.

- The DBMS better make sure one account deduction doesn't get lost..

Compare: An OS allows two people to edit a document at the same time. If both write, one's changes get lost.

Transactions

Transaction = process involving database queries and/or modification.

Normally with some strong properties regarding concurrency.

Formed in SQL from single statements.

ACID Transactions

ACID transactions are:

- *Atomic* : Whole transaction or none is done.
- *Consistent* : Database constraints preserved.
- *Isolated* : It appears to the user as if only one process executes at a time.
- *Durable* : Effects of a process survive a crash.

Optional: weaker forms of transactions are often supported as well.

COMMIT

The SQL statement **COMMIT** causes a **transaction to complete**.

- It's database modifications are now permanent in the database.

ROLLBACK

The SQL statement **ROLLBACK** also causes the transaction to end, but by *aborting*.

- No effects on the database.

Failures like division by 0 or a constraint violation can also cause rollback, even if the programmer does not request it.

Example: Interacting Processes

Assume the usual **Sells(bar,beer,price)** relation, and suppose that Joe's Bar sells only Bud for \$2.50 and Miller for \$3.00.

Sally is querying **Sells** for the highest (Query 1) and lowest (Query2) price Joe charges.

Joe decides to stop selling Bud and Miller (Query 3), but to sell only Heineken at \$3.50 (Query 4).

Sally's Program

Sally executes the following two SQL statements called (min) and (max) to help us remember what they do.

(max) SELECT MAX(price) FROM Sells
 WHERE bar = 'Joe''s Bar';

(min) SELECT MIN(price) FROM Sells
 WHERE bar = 'Joe''s Bar';

Joe's Program

At about the **same time**, Joe executes the following steps: (del) and (ins).

(del) DELETE FROM Sells
 WHERE bar = 'Joe''s Bar';

(ins) INSERT INTO Sells
 VALUES('Joe''s Bar', 'Heineken', 3.50);

Interleaving of Statements

Although (max) must come before (min), and (del) must come before (ins), there are **no other constraints on the order** of these statements, unless we group Sally's and/or Joe's statements into transactions!

Example: Strange Interleaving

Suppose the steps execute in the order `(max)(del)(ins)(min)`.

What is the result of Query 1 and Query 2?

Example: Strange Interleaving

Suppose the steps execute in the order (max)(del)(ins)(min).

Joe's Prices:	{2.50,3.00}	{2.50,3.00}		{3.50}
Statement:	(max)	(del)	(ins)	(min)
Result:	3.00			3.50

Sally sees MAX < MIN!

Fixing the Problem by Using Transactions

If we **group** Sally's statements **(max)(min)** into one transaction, then she cannot see this inconsistency.

She sees Joe's prices at some fixed time.

- Either before or after he changes prices, or in the middle, but the MAX and MIN are computed from the same prices.

Another Problem: Rollback

Suppose Joe executes **(del)(ins)**, **not as a transaction**, but after executing these statements, thinks better of it and issues a ROLLBACK statement.

If Sally executes her statements after **(ins)** but before the rollback, she sees a value, 3.50, that never existed in the database.

Solution

If Joe executes **(del)(ins)** as a **transaction**, its effect cannot be seen by others until the transaction executes COMMIT.

- If the transaction executes ROLLBACK instead, then its effects can *never* be seen.

Isolation Levels

SQL defines four *isolation levels* = choices about **what interactions are allowed by transactions** that execute at about the same time.

Only one level (“serializable”) = ACID transactions.

Each DBMS implements transactions in its own way.

Choosing the Isolation Level

Within a transaction, we can say:

SET TRANSACTION ISOLATION LEVEL X

where X =

1. SERIALIZABLE
2. REPEATABLE READ
3. READ COMMITTED
4. READ UNCOMMITTED

Serializable Transactions

If Sally = (max)(min) and Joe = (del)(ins) are each transactions, and **Sally** runs with **isolation level SERIALIZABLE**, then she will see the database either before or after Joe runs, but not in the middle.

Isolation Level Is Personal Choice

Your choice, e.g., run serializable, affects only how *you* see the database, not how others see it.

Example: If **Joe** Runs **serializable**, but Sally doesn't, then Sally might see no prices for Joe's Bar.

- i.e., it looks to Sally as if she ran in the middle of Joe's transaction.

Read-Committed Transactions

If **Sally** runs with isolation level **READ COMMITTED**, then she can see only committed data, but not necessarily the same data each time.

Example: Under READ COMMITTED, the interleaving **(max)(del)(ins)(min)** is allowed, as long as Joe commits.

- Sally sees $MAX < MIN$.

Repeatable-Read Transactions

Requirement is like read-committed, plus: if data is read again, then everything seen the first time will be seen the second time.

- But the second and subsequent reads may see *more* tuples as well.

Example: Repeatable Read

Suppose **Sally** runs under **REPEATABLE READ**, and the order of execution is (max)(del)(ins)(min).

- (max) sees prices 3.00.
- (min) can see 3.50, but must also see 3.00, because they were seen on the earlier read by (max).

Read Uncommitted

A transaction running under **READ UNCOMMITTED** can see data in the database, even if it was written by a transaction that has not committed (and may never).

Example: If Sally runs under READ UNCOMMITTED, she could see a price 3.50 even if Joe later aborts.

Views

A *view* is a relation defined in terms of stored tables (called *base tables*) and other views.

Two kinds:

1. *Virtual* = **not stored** in the database; just a query for constructing the relation.
2. *Materialized* = actually **constructed** and **stored**.

Declaring Views

Declare by:

```
CREATE [MATERIALIZED] VIEW    <name> AS <query>;
```

- Default is virtual.

Example: View Definition

CanDrink(drinker, beer) is a view “containing” the drinker-beer pairs such that the drinker frequents at least one bar that serves the beer:

Frequents(drinker, bar)

Sells(bar, beer, price)

```
CREATE VIEW CanDrink AS
  SELECT drinker, beer
  FROM Frequents, Sells
  WHERE Frequents.bar = Sells.bar;
```

Example: Accessing a View

Query a view as if it were a base table.

- Also: a limited ability to modify views if it makes sense as a modification of one underlying base table.

Example query:

```
SELECT beer FROM CanDrink  
WHERE drinker = 'Sally';
```

Triggers on Views

Generally, it is impossible to modify a virtual view, because it doesn't exist.

But an INSTEAD OF trigger lets us interpret view modifications in a way that makes sense.

Example: View **Synergy** has **(drinker, beer, bar)** triples such that the bar serves the beer, the drinker frequents the bar and likes the beer.

Example: The View

Pick one copy of
each attribute

CREATE VIEW Synergy AS

```
SELECT Likes.drinker, Likes.beer, Sells.bar
```

```
FROM Likes, Sells, Frequents
```

```
WHERE Likes.drinker = Frequents.drinker
```

```
    AND Likes.beer = Sells.beer
```

```
    AND Sells.bar = Frequents.bar;
```

Natural join of Likes,
Sells, and Frequents

Frequents(rinker, bar)
Sells(bar, beer, price)
Likes(rinker, beer)

Interpreting a View Insertion

We cannot insert into Synergy --- it is a virtual view.

But we can use an **INSTEAD OF trigger** to turn a (drinker, beer, bar) triple into three insertions of projected pairs, one for each of Likes, Sells, and Frequents.

- Sells.price will have to be NULL (as price is not an element of the view).

The Trigger

```
CREATE TRIGGER ViewTrig
  INSTEAD OF INSERT ON Synergy
  REFERENCING NEW ROW AS n
  FOR EACH ROW
  BEGIN
    INSERT INTO LIKES VALUES(n.drinker, n.beer);
    INSERT INTO SELLS(bar, beer) VALUES(n.bar, n.beer);
    INSERT INTO FREQUENTS VALUES(n.drinker, n.bar);
  END;
```


Materialized Views

Problem: each time a base table changes, the materialized view may change.

- Cannot afford to recompute the view with each change.

Solution: **Periodic reconstruction** of the **materialized view**, which is otherwise “out of date.”

Example: Class Mailing List

Eg., the class mailing list `cs-students` can be a **materialized view** of the class enrollment.

Actually updated four times/day.

- You can enroll and miss an email sent out after you enroll.

Example: A Data Warehouse

Wal-Mart stores every sale at every store in a database.

Overnight, the sales for the day are used to update a *data warehouse* = materialized views of the sales.

The warehouse is used by analysts to predict trends and move goods to where they are selling best.

Indexes

Index = **data structure** used to **speed access** to tuples of a relation, given values of one or more attributes.

Could be a **hash table**, but in a DBMS it is usually a **balanced search tree** with giant nodes (a full disk page) called a ***B-tree***.

Declaring Indexes

No standard!

Typical syntax:

```
CREATE INDEX BeerInd ON Beers (manf) ;
```

```
CREATE INDEX SellInd ON Sells (bar, beer) ;
```

Using Indexes

Given a value v , the index takes us to only those tuples that have v in the attribute(s) of the index.

Example: use BeerInd and SellInd to find the prices of beers manufactured by Pete's and sold by Joe. (next slide)

Using Indexes --- (2)

```
SELECT price FROM Beers, Sells
```

```
WHERE manf = 'Pete''s' AND
```

```
    Beers.name = Sells.beer AND
```

```
    bar = 'Joe''s Bar';
```

1. Use BeerInd to get all the beers made by Pete's.
2. Then use SellInd to get prices of those beers, with bar = 'Joe''s Bar'

Beers(name, manf)
Sells(bar, beer, price)

Database Tuning

A major problem in **making a database run fast** is **deciding which indexes to create**.

Pro: An index **speeds up queries** that can use it.

Con: An index **slows down all modifications** on its relation because the index must be modified too.

Example: Tuning

Suppose the only things we did with our beers database was:

1. Insert new facts into a relation (10%).
2. Find the price at a given bar of a given beer (90%).

Then **SellInd** on Sells(bar, beer) would be wonderful, but **BeerInd** on Beers(manf) would be harmful.

Tuning Advisors

A major research thrust.

- Because hand tuning is so hard.

An advisor gets a *query load*, e.g.:

1. Choose random queries from the history of queries run on the database, or
2. Designer provides a sample workload.

Tuning Advisors --- (2)

The **advisor** generates **candidate indexes** and evaluates each on the workload.

- Feed each sample query to the query optimizer, which assumes only this one index is available.
- Measure the improvement/degradation in the average running time of the queries.

Actions

Review slides!

Read Chapter 6.6 Transactions in SQL and Chapter 8 Views and Indexes.